

PROJECT TITLE AND WHAT IT DOES

Google Drive, read-only. Gives the AI safe, permission-scoped access to one real app so it can find and read my actual files — no copy-pasting, no uploading. The point of the task is less the summary it produces and more knowing exactly what access I handed over.

The one sentence the task asks for:
I granted read-only Google Drive access, letting Claude list and read my files but not create, edit, or delete them; I'd only need write access if I wanted it to save generated decks back to Drive rather than hand me the file to upload myself.

AI TOOLS AND APPS USED

TOOL / APP	ROLE
Claude (Opus 4.8)	The agent making the request
Google Drive connector	Read-only access — search files, read content. No write scope granted.

THE PROMPT I USED

Find my most recent document in my Drive, pull the three facts or numbers that matter most, and give me a one-line summary of what it is.

Deliberately vague about *which* file. Naming the file would have proved nothing — I wanted the connector to do the finding, not just the fetching.

WHAT CAME BACK

Claude found `skills-connectors.pptx` in my *Agent Factory Decks* folder, modified 14 July.

One-line summary: a slide deck of the Skills & Connectors crash course, built from the Agent Factory book.

The three that mattered:

- Structure: 5 parts + 1 core idea.** The deck's spine — a chat message tells AI what to do *once*, a Skill teaches it *how*, *every time*, a Connector gives it *hands*.
- Positioning: "2 features. 5 tools. 0 code from you."** The subtitle, and the promise the deck is built to keep.
- Skills fire automatically ~90% of the time.** You describe the task; AI scans it against every skill's description; the match loads on its own.

HOW I TESTED AND VERIFIED IT

Opened the real `skills-connectors.pptx` in Drive and checked each claim against the actual slides:

CLAIM	CHECKED AGAINST
Title and folder location	File sits in <i>Agent Factory Decks</i> , modified 14 July
"5 parts + 1 core idea"	Slide 2 — the contents slide lists exactly five parts
"2 features. 5 tools. 0 code from you."	Slide 1 subtitle, word for word
"~90% automatic"	The slide on invocation

This step is the entire point of the task. A connector returning a *confident* summary and one returning a *correct* summary look identical from the chat window. The only way to tell them apart is to open the source — which is why the course rule is "always check the AI's output against the real source before you trust it."

WOULD I EVER NEED WRITE ACCESS?

Yes, for one specific workflow. My *Agent Factory Decks* folder is one I own and can add to — so if I wanted Claude to generate a deck and **save it back to Drive** rather than hand me a file to upload myself, that is a write action and read-only would not cover it.

The tradeoff I would be accepting: write access means a wrong edit, a file created in the wrong place, or a deletion with no easy undo. Read-only has no such failure mode — the worst case is a bad summary, which I catch by checking the source. **Staying read-only until the copy-pasting actually hurts is the cheaper default.**

WHAT WORKED, WHAT DID NOT, AND PROBLEMS I FACED

What worked

- The vague request was handled correctly.** "My most recent document" required the connector to search and sort, not just fetch a named file.
- Read-only was genuinely sufficient.** Nothing in this workflow needed write access, which is the course's point about scope: connect what you need, nothing more.
- AI inherits my permissions.** The connector could only reach what I can already reach — connecting Drive did not expose anything I don't already own.

What did not work / problems faced

- The verification step nearly got skipped.** The summary was fluent, correctly formatted, and matched what I expected the deck to say — which is exactly the condition under which people stop checking. I had to deliberately open the file to confirm, because nothing about the output signalled a need to.
- The result exposed my email address.** The connector returns file owner metadata, which put my Gmail address in the output. Not a failure — but it is a reminder that connector output carries more than the content you asked for, and it has to be blurred before anything is committed to a public repo.

The lesson I actually took: the permission screen is the task. Clicking Connect takes two seconds and reading what you are agreeing to takes twenty — and the twenty seconds are the difference between "I connected Drive" and "I know what Drive access means." That distinction is what Task 5 grades.

